

Loading the data from cloud

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: import pandas as pd

df=pd.read_csv('/content/drive/MyDrive/CarPrice_Assignment.csv')
df.shape
```

Out[]: (205, 26)

```
In [ ]: df.head()
```

Out[]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	eng
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	

5 rows × 26 columns

◀

```
In [ ]: df.info()
```

▶

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration            205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel           205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength             205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight             205 non-null    float64
13  curbweight            205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize            205 non-null    int64
17  fuelsystem            205 non-null    object
18  boreratio             205 non-null    float64
19  stroke                205 non-null    float64
20  compressionratio      205 non-null    float64
21  horsepower            205 non-null    int64
22  peakrpm              205 non-null    int64
23  citympg              205 non-null    int64
24  highwaympg           205 non-null    int64
25  price                205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
In [ ]: df.describe()
```

```
Out [ ]:
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	engin
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.00
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.90
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.64
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.00
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.00
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.00
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.00
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.00

```

In [ ]: #Splitting company name from CarName column
CompanyName = df['CarName'].apply(lambda x : x.split(' ')[0])
df.insert(3,"CompanyName",CompanyName)
df.drop(['CarName'],axis=1,inplace=True)
df.head()

```

```
Out[ ]:
```

	car_ID	symboling	CompanyName	fueltype	aspiration	doornumber	carbody	drivewheel
0	1	3	alfa-romero	gas	std	two	convertible	rwd
1	2	3	alfa-romero	gas	std	two	convertible	rwd
2	3	1	alfa-romero	gas	std	two	hatchback	rwd
3	4	2	audi	gas	std	four	sedan	fwd
4	5	2	audi	gas	std	four	sedan	4wd

5 rows × 9 columns

```
In [ ]:
```

```
df.CompanyName = df.CompanyName.str.lower()

def replace_name(a,b):
    df.CompanyName.replace(a,b,inplace=True)

replace_name('maxda','mazda')
replace_name('porcshce','porsche')
replace_name('toyouta','toyota')
replace_name('vokswagen','volkswagen')
replace_name('vw','volkswagen')

print(df.CompanyName.unique())
df.CompanyName.value_counts()
```

```
Out[ ]:
```

```
['alfa-romero' 'audi' 'bmw' 'chevrolet' 'dodge' 'honda' 'isuzu' 'jaguar'
 'mazda' 'buick' 'mercury' 'mitsubishi' 'nissan' 'peugeot' 'plymouth'
 'porsche' 'renault' 'saab' 'subaru' 'toyota' 'volkswagen' 'volvo']
```

```
toyota      32
nissan       18
mazda       17
mitsubishi  13
honda       13
volkswagen  12
subaru      12
peugeot     11
volvo       11
dodge        9
buick        8
bmw          8
audi         7
plymouth     7
saab         6
porsche      5
isuzu        4
jaguar       3
chevrolet    3
alfa-romero  3
renault      2
mercury      1
Name: CompanyName, dtype: int64
```

Exploratory data analysis

Histogram analysis for numerical variables

```
In [ ]: #Histogram for car price
import matplotlib.pyplot as plt
import seaborn as sns

sns.distplot(df.price)
```

<ipython-input-8-3501d61547f0>:5: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.price)
```

Out[]: <Axes: xlabel='price', ylabel='Density'>



```
In [ ]: sns.distplot(df.horsepower)
```

<ipython-input-18-fe36dc2a0d7c>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.horsepower)
```


Out[]: <Axes: xlabel='horsepower', ylabel='Density'>



Barplot analysis for categorical variables


```
In [ ]: sns.barplot(data=df,x='carbody',y='price')
```

```
Out[ ]: <Axes: xlabel='carbody', ylabel='price'>
```




```
In [ ]: sns.barplot(data=df,x='drivewheel',y='price')
```

```
Out[ ]: <Axes: xlabel='drivewheel', ylabel='price'>
```




```
In [ ]: sns.barplot(data=df,x='fuelsystem',y='price')
```

```
Out[ ]: <Axes: xlabel='fuelsystem', ylabel='price'>
```




```
In [ ]: sns.barplot(data=df,y='CompanyName',x='price',orient='h')
```

```
Out[ ]: <Axes: xlabel='price', ylabel='CompanyName'>
```



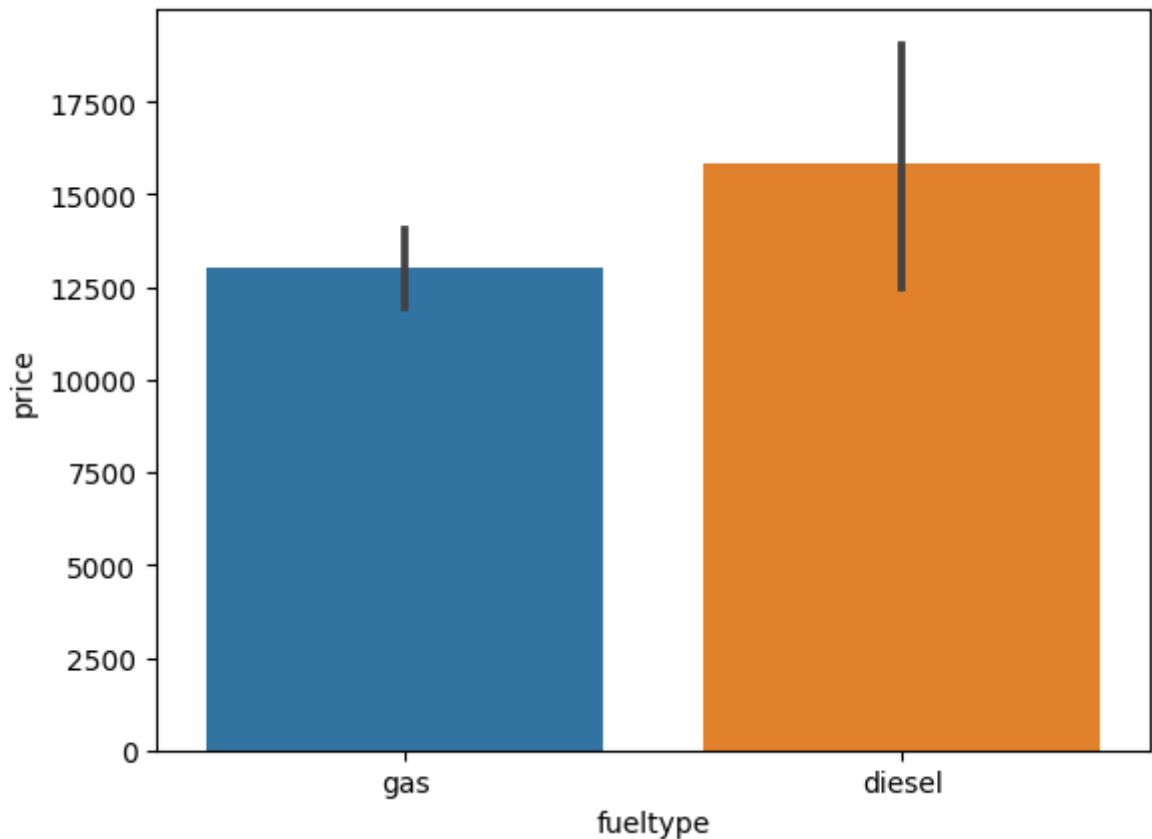
```
In [ ]: plt01 = df.CompanyName.value_counts().plot(kind='bar')
plt.title('Companies Histogram')
plt01.set(xlabel = 'Car company', ylabel='Frequency of company')
```

```
Out[ ]: [Text(0.5, 0, 'Car company'), Text(0, 0.5, 'Frequency of company')]
```



```
In [ ]: sns.barplot(data=df,x='fueltype',y='price')
```

```
Out[ ]: <Axes: xlabel='fueltype', ylabel='price'>
```




Correlation analysis using heatmap

```
In [ ]: corr = df.corr()  
plt.figure(figsize=(12, 9))  
sns.heatmap(corr, annot = True)
```

<ipython-input-17-ca7bf234e79c>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr = df.corr()
```

```
Out[ ]: <Axes: >
```

```
In [ ]: selected_cols=df[['wheelbase','carlength','carwidth','carheight','curbweight','enginesize','citympg','highwaympg','price']]
#sns.pairplot(selected_cols)
```

Checking multicollinearity using VIF

```
In [ ]: #Defining model and checking multicollinearity of variables
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
def build_model(X,y):
    X = sm.add_constant(X) #Adding the constant
    lm = sm.OLS(y,X).fit() # fitting the model
    print(lm.summary()) # model summary
    return X
def checkVIF(X):
    vif = pd.DataFrame()
    vif['Features'] = X.columns
    vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)
    return(vif)
checkVIF(selected_cols)
```

Out[]:

	Features	VIF
2	carwidth	2350.90
0	wheelbase	1903.66
1	carlength	1893.74
3	carheight	917.20
12	highwaympg	508.20
11	citympg	429.21
4	curbweight	403.02
6	boreratio	290.28
10	peakrpm	217.66
7	stroke	125.83
5	enginesize	68.77
9	horsepower	65.53
8	compressionratio	15.75

In []:

```
#Dividing data into X and y variables  
from sklearn.model_selection import train_test_split  
y = selected_cols.pop('price')  
X = selected_cols  
X_train, X_test, y_train, y_test = train_test_split(X,y , random_state=104, train_s
```

In []:

```
#Model1  
md1=build_model(X_train,y_train)
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.870
Model:                  OLS        Adj. R-squared:           0.859
Method:                 Least Squares  F-statistic:             77.19
Date:                  Wed, 13 Sep 2023  Prob (F-statistic):      1.52e-59
Time:                  17:08:29     Log-Likelihood:          -1540.5
No. Observations:      164         AIC:                    3109.
Df Residuals:          150         BIC:                    3152.
Df Model:              13
Covariance Type:       nonrobust
=====
==

```

	coef	std err	t	P> t	[0.025	0.975]
const	-3.894e+04	1.6e+04	-2.441	0.016	-7.05e+04	-7420.1
wheelbase	105.2326	105.964	0.993	0.322	-104.141	314.6
carlength	-64.6334	59.625	-1.084	0.280	-182.446	53.1
carwidth	412.0404	245.177	1.681	0.095	-72.406	896.4
carheight	214.9141	149.468	1.438	0.153	-80.421	510.2
curbweight	0.3725	1.820	0.205	0.838	-3.223	3.9
enginesize	152.3401	15.820	9.629	0.000	121.081	183.6
boreratio	-2627.9343	1382.323	-1.901	0.059	-5359.274	103.4
stroke	-4037.9994	904.330	-4.465	0.000	-5824.869	-2251.1
compressionratio	323.4195	87.106	3.713	0.000	151.305	495.5
horsepower	32.3381	16.266	1.988	0.049	0.197	64.4
peakrpm	2.3312	0.693	3.365	0.001	0.962	3.7
citympg	-330.7249	184.534	-1.792	0.075	-695.347	33.8
highwaympg	189.2060	161.815	1.169	0.244	-130.524	508.9

```

=====
Omnibus:                17.839    Durbin-Watson:           2.084
Prob(Omnibus):           0.000    Jarque-Bera (JB):        70.158
Skew:                   0.068    Prob(JB):                5.83e-16
Kurtosis:               6.201    Cond. No.                 3.87e+05
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.87e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In []: `checkVIF(md1)`

Out[]:

	Features	VIF
0	const	4524.13
12	citympg	24.15
13	highwaympg	20.92
5	curbweight	15.42
2	carlength	9.44
10	horsepower	7.21
1	wheelbase	7.19
6	enginesize	6.71
3	carwidth	4.84
7	boreratio	2.47
4	carheight	2.29
9	compressionratio	2.24
11	peakrpm	2.04
8	stroke	1.43

Feature engineering

In []:

```
#Dropping the columns with high multicollinearity  
X_train_2 = X_train.drop(['citympg', 'highwaympg', 'curbweight'], axis=1)
```

In []:

```
#Model2  
md2=build_model(X_train_2,y_train)  
checkVIF(md2)
```

OLS Regression Results

```

=====
Dep. Variable:          price      R-squared:                0.866
Model:                  OLS        Adj. R-squared:           0.858
Method:                 Least Squares    F-statistic:             99.15
Date:                  Wed, 13 Sep 2023    Prob (F-statistic):      1.28e-61
Time:                  17:08:44      Log-Likelihood:          -1542.7
No. Observations:      164          AIC:                    3107.
Df Residuals:          153          BIC:                    3142.
Df Model:              10
Covariance Type:       nonrobust
=====

```

```

=====
==
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
--
const          -4.827e+04   1.43e+04    -3.367     0.001   -7.66e+04   -1.99e+
04
wheelbase       75.7243     101.550     0.746     0.457   -124.896    276.3
45
carlength      -16.2208      53.224    -0.305     0.761   -121.370     88.9
28
carwidth        421.2682    243.569     1.730     0.086    -59.925    902.4
61
carheight       193.3858    149.489     1.294     0.198   -101.942    488.7
14
enginesize      149.8949     14.264    10.509     0.000    121.716    178.0
74
boreratio      -2226.0409   1371.615    -1.623     0.107   -4935.791    483.7
09
stroke          -3887.6179    900.697    -4.316     0.000   -5667.026   -2108.2
10
compressionratio  269.7576     71.262     3.785     0.000    128.973    410.5
42
horsepower      46.2314     14.465     3.196     0.002     17.654     74.8
09
peakrpm         2.4048       0.684     3.517     0.001     1.054       3.7
56
=====

```

```

=====
Omnibus:          16.333    Durbin-Watson:           2.113
Prob(Omnibus):    0.000    Jarque-Bera (JB):        56.493
Skew:             0.114    Prob(JB):                5.40e-13
Kurtosis:         5.866    Cond. No.                 3.10e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.1e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Out[]:

	Features	VIF
0	const	3625.19
2	carlength	7.46
1	wheelbase	6.55
9	horsepower	5.66
5	enginesize	5.41
3	carwidth	4.74
6	boreratio	2.41
4	carheight	2.27
10	peakrpm	1.97
8	compressionratio	1.49
7	stroke	1.40

```
In [ ]: X_train_3 =X_train_2.drop(['carlength'],axis=1)
```

```
In [ ]: #Model3
md3=build_model(X_train_3,y_train)
checkVIF(md3)
```

OLS Regression Results

```

=====
Dep. Variable:          price      R-squared:                0.866
Model:                  OLS        Adj. R-squared:            0.858
Method:                 Least Squares    F-statistic:             110.8
Date:                   Wed, 13 Sep 2023    Prob (F-statistic):      1.24e-62
Time:                   17:09:04      Log-Likelihood:          -1542.8
No. Observations:      164          AIC:                     3106.
Df Residuals:          154          BIC:                     3137.
Df Model:               9
Covariance Type:       nonrobust
=====

```

```

=====
==
              coef      std err          t      P>|t|      [0.025      0.975
-----
5]
-----
--
const      -4.701e+04    1.37e+04     -3.433     0.001    -7.41e+04    -2e+
04
wheelbase      58.9255      85.037       0.693     0.489    -109.063     226.9
14
carwidth      404.4800     236.558       1.710     0.089    -62.838     871.7
98
carheight      179.2757     141.719       1.265     0.208    -100.689     459.2
40
enginesize      149.6247      14.194      10.541     0.000     121.584     177.6
65
boreratio     -2320.6541    1332.080      -1.742     0.083   -4952.162     310.8
54
stroke        -3923.4789     890.344      -4.407     0.000   -5682.343   -2164.6
15
compressionratio  271.3247      70.867       3.829     0.000     131.329     411.3
21
horsepower      45.5711      14.260       3.196     0.002      17.401      73.7
41
peakrpm         2.3988       0.681       3.520     0.001       1.053       3.7
45
=====

```

```

=====
Omnibus:          15.997    Durbin-Watson:           2.107
Prob(Omnibus):    0.000    Jarque-Bera (JB):        54.874
Skew:             0.092    Prob(JB):                1.21e-12
Kurtosis:         5.828    Cond. No.                 2.97e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.97e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Out[]:

	Features	VIF
0	const	3327.46
8	horsepower	5.53
4	enginesize	5.39
1	wheelbase	4.62
2	carwidth	4.50
5	boreratio	2.29
3	carheight	2.05
9	peakrpm	1.97
7	compressionratio	1.48
6	stroke	1.38

Elimination based on multicollinearity,p-value and the correlation

```
In [ ]: #Dropping the columns with high p-value
X_train_4 = X_train_3.drop(['wheelbase','carheight','boreratio'],axis=1)
```

```
In [ ]: md4=build_model(X_train_4,y_train)
checkVIF(md4)
```


Dep. Variable:	price	R-squared:	0.860
Model:	OLS	Adj. R-squared:	0.855
Method:	Least Squares	F-statistic:	161.3
Date:	Wed, 13 Sep 2023	Prob (F-statistic):	1.75e-64
Time:	17:09:18	Log-Likelihood:	-1546.3
No. Observations:	164	AIC:	3107.
Df Residuals:	157	BIC:	3128.
Df Model:	6		
Covariance Type:	nonrobust		

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.42e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [ ]: col=X_train_4.columns
X_test_1=X_test[['carwidth', 'enginesize', 'stroke', 'compressionratio', 'horsepower']]

In [ ]: # Predict using the trained model on the test set
# Making predictions
```

```

from sklearn.linear_model import LinearRegression
lm = LinearRegression()
fitted=lm.fit(X_train_4,y_train)
y_pred = lm.predict(X_test_1)

# You can also calculate metrics to evaluate the predictions, for example, Mean Squared Error
from sklearn.metrics import mean_squared_error

mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)

```

Mean Squared Error: 17045747.076587398

```

In [ ]: from sklearn.metrics import r2_score
        r2_score(y_test, y_pred)

```

Out[]: 0.704915490316405

```


In [ ]: # Compute residuals
        residuals = y_test - y_pred

# Compute Mean Squared Error (MSE) to assess model fit
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")

# Plot residuals vs. predicted values
plt.scatter(y_pred, residuals)
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residuals vs. Predicted Values")
plt.show()

```

Mean Squared Error: 17045747.076587398

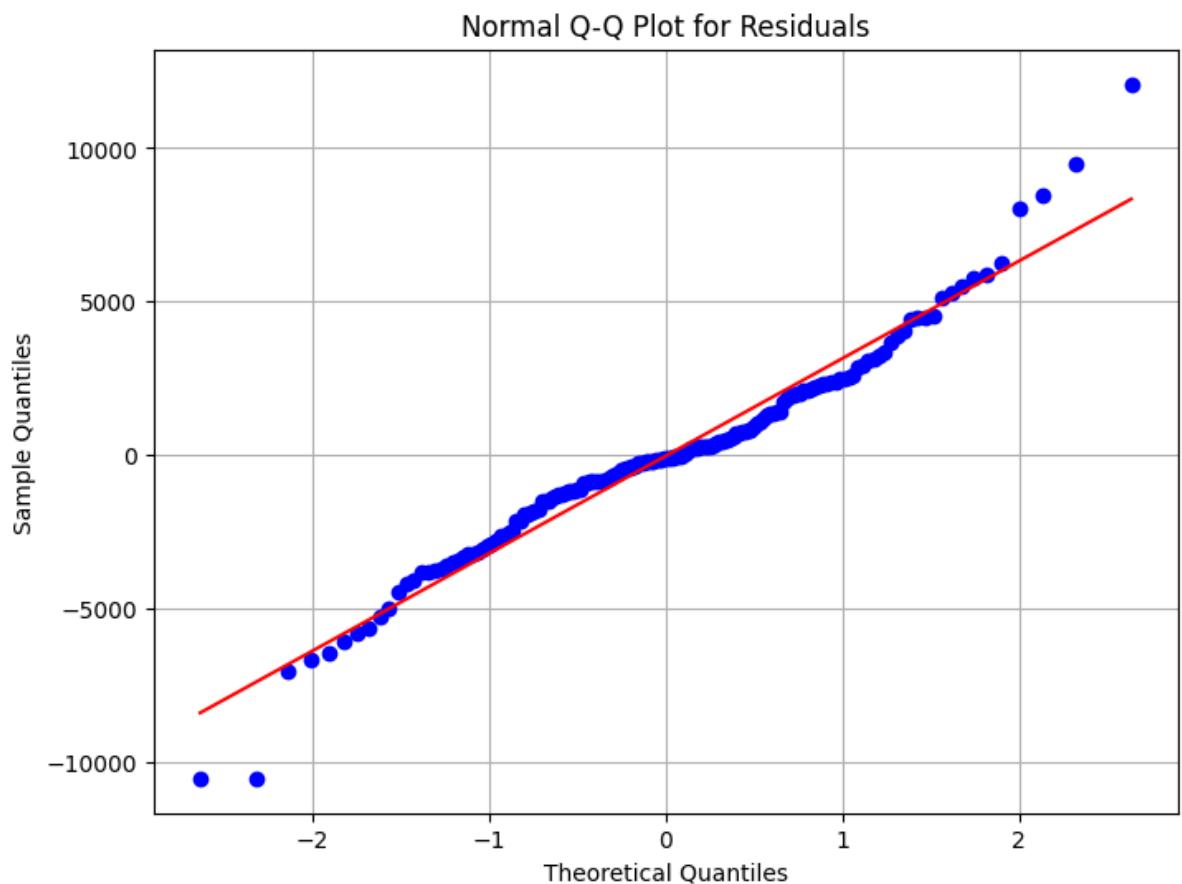


```
In [ ]: import statsmodels.api as sm
import scipy.stats as stats

model = sm.OLS(y_train, X_train_4).fit()

# Get the residuals from the model
residuals = model.resid

# Create a Normal Q-Q plot for the residuals
plt.figure(figsize=(8, 6))
stats.probplot(residuals, dist="norm", plot=plt)
plt.title("Normal Q-Q Plot for Residuals")
plt.xlabel("Theoretical Quantiles")
plt.ylabel("Sample Quantiles")
plt.grid(True)
plt.show()
```



```
In [ ]: import numpy as np
user_input = []
for col in range(X_train_4.shape[1]): # Iterate over the number of columns
    value = float(input(f"Enter value for {X_train_4.columns[col]}: "))
    user_input.append(value)

# Predict using the trained model
user_input_array = np.array(user_input).reshape(1, -1) # Reshape for prediction
user_input_array_with_constant = sm.add_constant(user_input_array) # Add constant
predicted_value = model.predict(user_input_array_with_constant)

print("Predicted car price:", predicted_value[0])
```

```
Enter value for carwidth: 64
Enter value for enginesize: 130
Enter value for stroke: 2.2
Enter value for compressionratio: 8
Enter value for horsepower: 100
Enter value for peakrpm: 4500
Predicted car price: 16873.61509462775
```

In []: